



Matched-three-run sensitivity across 15 SR methods

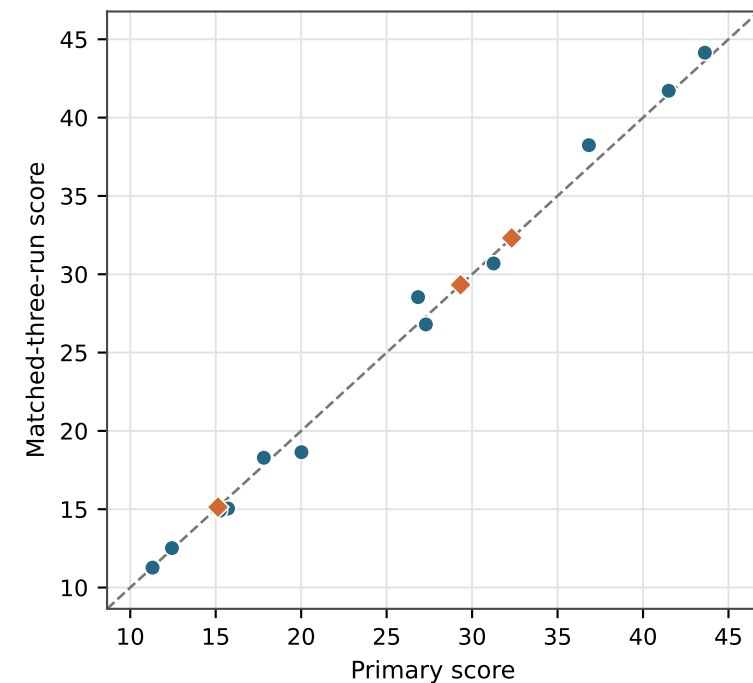
ID-Q (rank $r=0.993$)



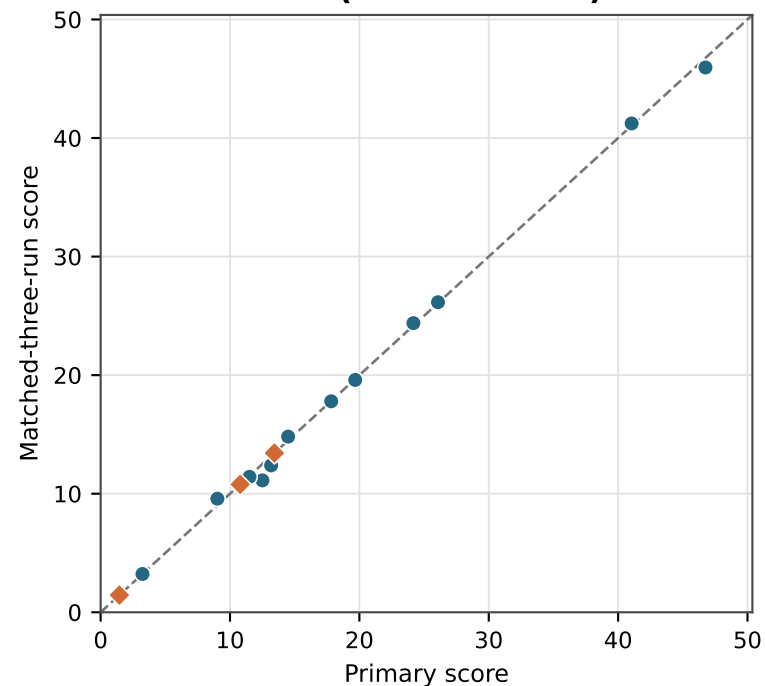
OOD-G (rank $r=0.979$)



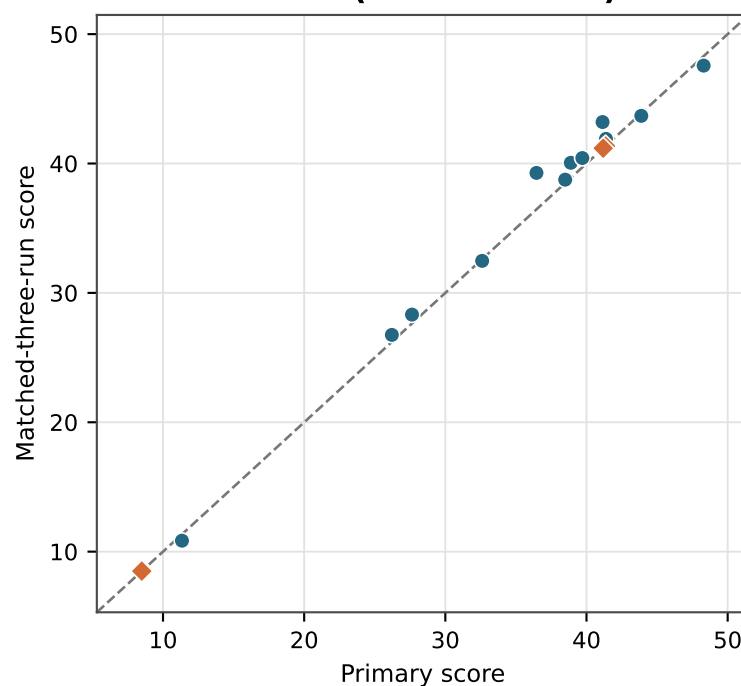
SYM-F (rank $r=0.986$)



EFF (rank $r=0.996$)



ROBU (rank $r=0.975$)



● Original 12
◆ Rebuttal-added 3

All methods use three runs per task; ROBU uses the two common noise levels. STAB is excluded because its matched seed-level proxy is not part of this sensitivity artifact.